



Agent context lives everywhere.





Don't assemble your agents' context layer.
Use **Redis Iris**.

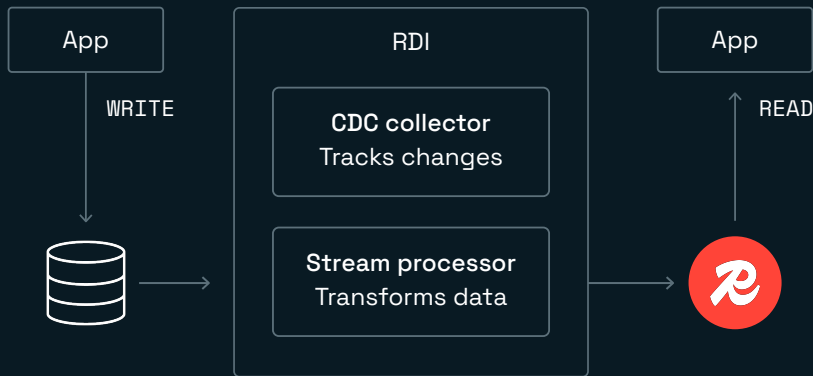
Powered by the Redis you already love



Redis Data Integration (RDI) makes your data fresh.

Keep your context in Redis synced from other data sources.

- Effortlessly configure and deploy a data pipeline between Redis and your other slower databases.
- Capture, ingest, and transform changed data with configuration—not code.
- Easily deploy, manage, and visualize your data pipeline in Redis Insight or the Redis Cloud console.



How RDI works:

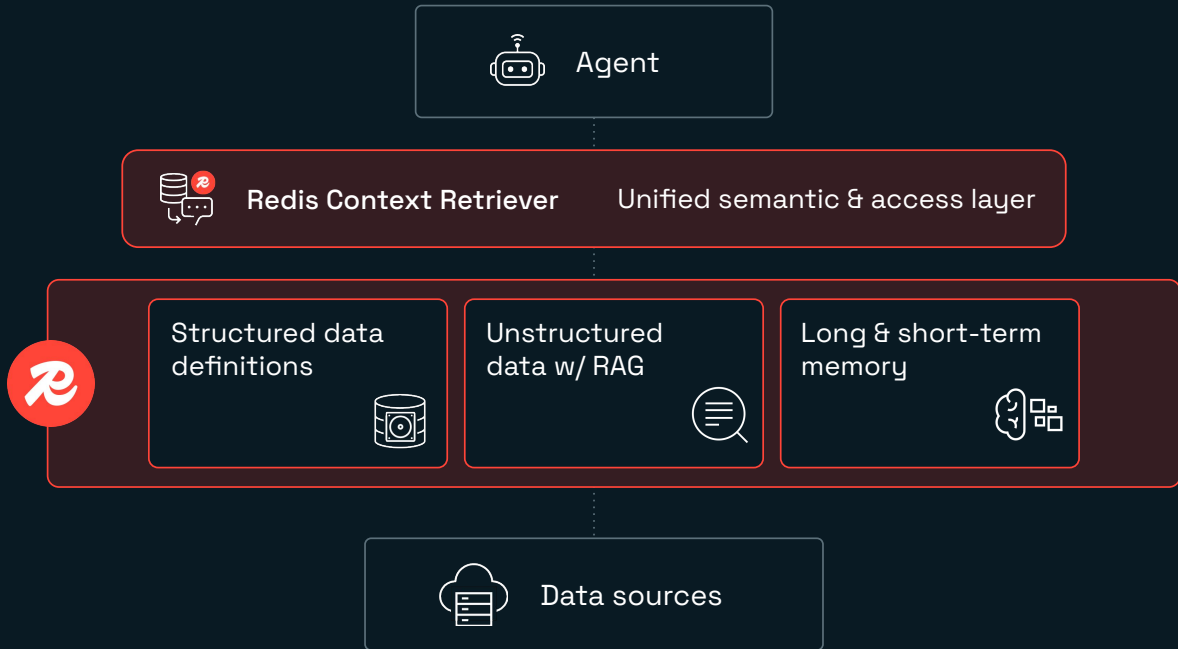
1. RDI initially populates Redis with data from your DB
2. A change data capture (CDC) collector keeps track of any changed data written to your database
3. RDI stream processor translates the data to the preferred data model and types
4. The data is ingested into Redis and can be read from your now faster app

Context Retriever makes your data navigable.

Turn your enterprise data into a surface agents can use with a unified semantic & access layer.

The context portal:

- **User:**
“Why’s my order late?”
- **Response:**
Hello Simba, let me look up your order ... I see! Your courier had a flat tire. Let me look up our policies ... I’ll provide you a refund, and get you a new deliverer!

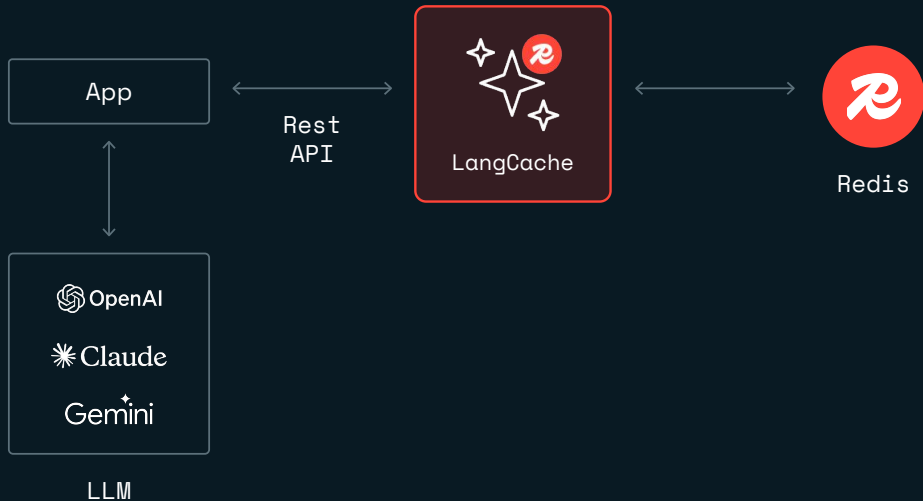


LangCache makes retrieval fast.

Improve agent results
with accelerated retrieval
using fewer tokens.

Semantic caching eliminates
redundant LLM calls.

- 15X faster
- 70% fewer tokens
- Highly accurate



Agent Memory makes your context compounding.

Store history and preferences, and retrieve key details for agents and customers.

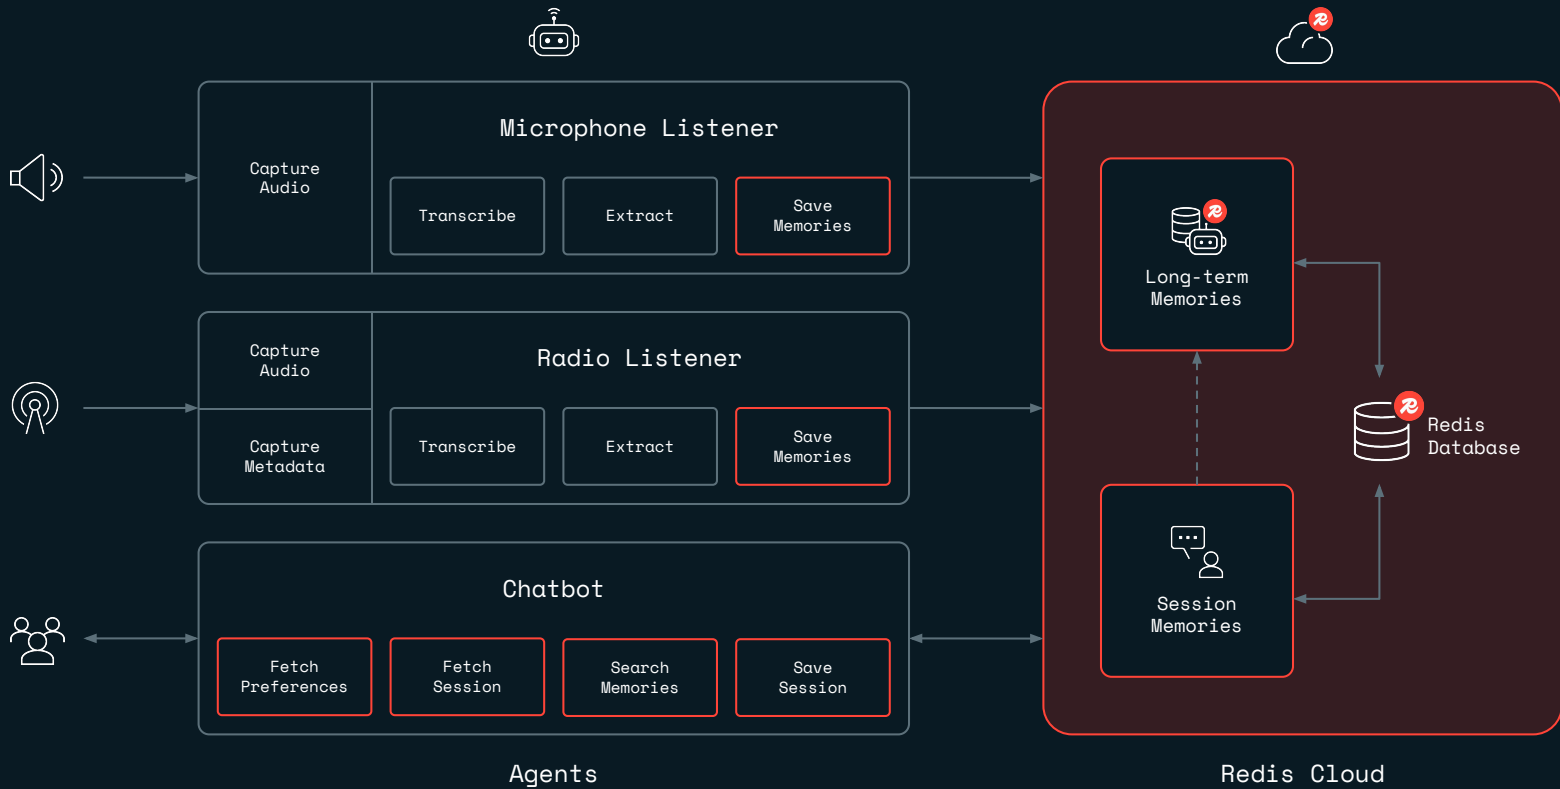
Short-term memory is typically isolated per user and per session.

- Agent state
- Chat history
- Summaries

Long-term memory is typically shared across sessions and sometimes across users.

- Extracted preferences
- Durable facts





Earshot Architecture

Demo